

MA415 Assignment 4

Kevin Chan

February 21 2018

10.5 Exercises

1.

How can you tell if an object is a tibble? (Hint: try printing `mtcars`, which is a regular data frame.)

```
# Initial setup.
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 3.4.3
```

```
## -- Attaching packages ----- tidyverse 1.2.1 --
```

```
## v ggplot2 2.2.1      v purrr  0.2.4
```

```
## v tibble  1.4.2      v dplyr  0.7.4
```

```
## v tidyr   0.7.2      v stringr 1.2.0
```

```
## v readr   1.1.1      v forcats 0.2.0
```

```
## Warning: package 'ggplot2' was built under R version 3.4.3
```

```
## Warning: package 'tibble' was built under R version 3.4.3
```

```
## Warning: package 'tidyr' was built under R version 3.4.3
```

```
## Warning: package 'readr' was built under R version 3.4.2
```

```
## Warning: package 'purrr' was built under R version 3.4.3
```

```
## Warning: package 'dplyr' was built under R version 3.4.3
```

```
## Warning: package 'stringr' was built under R version 3.4.3
```

```
## Warning: package 'forcats' was built under R version 3.4.3
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
library(knitr)
```

```
## Warning: package 'knitr' was built under R version 3.4.3
```

```
# Exercise 1.
```

```
mtcars
```

```
##           mpg  cyl  disp  hp drat   wt  qsec vs  am gear carb
## Mazda RX4      21.0    6 160.0 110 3.90 2.620 16.46 0   1    4    4
## Mazda RX4 Wag  21.0    6 160.0 110 3.90 2.875 17.02 0   1    4    4
## Datsun 710     22.8    4 108.0  93 3.85 2.320 18.61 1   1    4    1
## Hornet 4 Drive  21.4    6 258.0 110 3.08 3.215 19.44 1   0    3    1
## Hornet Sportabout 18.7    8 360.0 175 3.15 3.440 17.02 0   0    3    2
## Valiant        18.1    6 225.0 105 2.76 3.460 20.22 1   0    3    1
## Duster 360     14.3    8 360.0 245 3.21 3.570 15.84 0   0    3    4
## Merc 240D      24.4    4 146.7  62 3.69 3.190 20.00 1   0    4    2
```

```
## Merc 230      22.8  4 140.8  95 3.92 3.150 22.90  1 0   4   2
## Merc 280      19.2  6 167.6 123 3.92 3.440 18.30  1 0   4   4
## Merc 280C     17.8  6 167.6 123 3.92 3.440 18.90  1 0   4   4
## Merc 450SE    16.4  8 275.8 180 3.07 4.070 17.40  0 0   3   3
## Merc 450SL    17.3  8 275.8 180 3.07 3.730 17.60  0 0   3   3
## Merc 450SLC   15.2  8 275.8 180 3.07 3.780 18.00  0 0   3   3
## Cadillac Fleetwood 10.4  8 472.0 205 2.93 5.250 17.98  0 0   3   4
## Lincoln Continental 10.4  8 460.0 215 3.00 5.424 17.82  0 0   3   4
## Chrysler Imperial 14.7  8 440.0 230 3.23 5.345 17.42  0 0   3   4
## Fiat 128      32.4  4  78.7  66 4.08 2.200 19.47  1 1   4   1
## Honda Civic   30.4  4  75.7  52 4.93 1.615 18.52  1 1   4   2
## Toyota Corolla 33.9  4  71.1  65 4.22 1.835 19.90  1 1   4   1
## Toyota Corona 21.5  4 120.1  97 3.70 2.465 20.01  1 0   3   1
## Dodge Challenger 15.5  8 318.0 150 2.76 3.520 16.87  0 0   3   2
## AMC Javelin   15.2  8 304.0 150 3.15 3.435 17.30  0 0   3   2
## Camaro Z28    13.3  8 350.0 245 3.73 3.840 15.41  0 0   3   4
## Pontiac Firebird 19.2  8 400.0 175 3.08 3.845 17.05  0 0   3   2
## Fiat X1-9     27.3  4  79.0  66 4.08 1.935 18.90  1 1   4   1
## Porsche 914-2 26.0  4 120.3  91 4.43 2.140 16.70  0 1   5   2
## Lotus Europa  30.4  4  95.1 113 3.77 1.513 16.90  1 1   5   2
## Ford Pantera L 15.8  8 351.0 264 4.22 3.170 14.50  0 1   5   4
## Ferrari Dino   19.7  6 145.0 175 3.62 2.770 15.50  0 1   5   6
## Maserati Bora  15.0  8 301.0 335 3.54 3.570 14.60  0 1   5   8
## Volvo 142E    21.4  4 121.0 109 4.11 2.780 18.60  1 1   4   2
```

```
class(mtcars)
```

```
## [1] "data.frame"
```

```
as_tibble(mtcars)
```

```
## # A tibble: 32 x 11
##   mpg   cyl  disp    hp  drat    wt  qsec    vs  am  gear  carb
##   * <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 21.0   6.00  160 110   3.90  2.62  16.5   0     1.00  4.00  4.00
## 2 21.0   6.00  160 110   3.90  2.88  17.0   0     1.00  4.00  4.00
## 3 22.8   4.00  108  93.0  3.85  2.32  18.6   1.00  1.00  4.00  1.00
## 4 21.4   6.00  258 110   3.08  3.22  19.4   1.00  0     3.00  1.00
## 5 18.7   8.00  360 175   3.15  3.44  17.0   0     0     3.00  2.00
## 6 18.1   6.00  225 105   2.76  3.46  20.2   1.00  0     3.00  1.00
## 7 14.3   8.00  360 245   3.21  3.57  15.8   0     0     3.00  4.00
## 8 24.4   4.00  147  62.0  3.69  3.19  20.0   1.00  0     4.00  2.00
## 9 22.8   4.00  141  95.0  3.92  3.15  22.9   1.00  0     4.00  2.00
## 10 19.2   6.00  168 123   3.92  3.44  18.3   1.00  0     4.00  4.00
## # ... with 22 more rows
```

```
class(as_tibble(mtcars))
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

2.

Compare and contrast the following operations on a `data.frame` and equivalent tibble. What is different? Why might the default data frame behaviours cause your frustration? Using “\$” with a data frame will partially complete the column, which may result in accidentally using a different variable than the one

desired. Using `[]`, the type of object returned is dependent on the number of columns. If there is only one column, the object will return a vector and not a data frame. Otherwise, the object will return a data frame.

```
df <- data.frame(abc = 1, xyz = "a")
df$x
```

```
## [1] a
## Levels: a
```

```
df[, "xyz"]
```

```
## [1] a
## Levels: a
```

```
df[, c("abc", "xyz")]
```

```
##   abc xyz
## 1   1   a
```

```
tibdf <- as_tibble(df)
tibdf$x
```

```
## Warning: Unknown or uninitialised column: 'x'.
```

```
## NULL
```

```
tibdf[, "xyz"]
```

```
## # A tibble: 1 x 1
##   xyz
##   <fct>
## 1 a
```

```
tibdf[, c("abc", "xyz")]
```

```
## # A tibble: 1 x 2
##   abc xyz
##   <dbl> <fct>
## 1  1.00 a
```

3.

If you have the name of a variable stored in an object, e.g. `var <- "mpg"`, how can you extract the reference variable from a tibble?

To extract the reference variable from a tibble, you can use the double bracket, e.g. `df[[var]]`.

4.

Practice referring to non-syntactic names in the following data frame by:

1.

Extracting the variable called `\texttt{1}`.

```
# This is given.
annoying <- tibble(
  `1` = 1:10,
```

```

`2` = `1` * 2 + rnorm(length(`1`))
)
# Part 1 of Problem 4
annoying[["1"]]


## [1] 1 2 3 4 5 6 7 8 9 10

```

2.

Plotting a scatterplot of `\texttt{1}` vs `\texttt{2}`.

```
ggplot(annoying, aes(x = `1`, y = `2`)) + geom_point()
```



3.

Creating a new column called `\texttt{3}`, which is `\texttt{2}` divided by `\texttt{1}`.

```
annoying[["3"]] <- annoying[["2"]]/annoying[["1"]]
```

4.

Renaming the columns to `\texttt{one}`, `\texttt{two}`, and `\texttt{three}`.

```
annoying <- rename(annoying, one = `1`, two = `2`, three = `3`)
glimpse(annoying)
```

```
## Observations: 10
## Variables: 3
## $ one <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
## $ two <dbl> 1.871761, 5.477952, 4.832513, 6.991914, 8.057822, 11.564...
## $ three <dbl> 1.871761, 2.738976, 1.610838, 1.747978, 1.611564, 1.9274...
```

5.

What does `tibble::enframe()` do? When might you use it?

```
?tibble::enframe()
```

```
## starting httpd help server ... done
```

`tibble::enframe()` converts named vectors to two-column data frames with names and values. The natural sequence is used as name column for unnamed vectors. ## 6. What option controls how many additional column names are printed at the footer of a tibble? Using the print function for tibbles, which is `print.tbl_df`, the option `n_extra` controls how many extra column names are printed.

12.6.1 Exercises

```
# Code necessary for the exercises below (from throughout the chapter):
```

```
who %>%
  gather(code, value, new_sp_m014:newrel_f65, na.rm = TRUE) %>%
  mutate(code = stringr::str_replace(code, "newrel", "new_rel")) %>%
  separate(code, c("new", "var", "sexage")) %>%
  select(-new, -iso2, -iso3) %>%
  separate(sexage, c("sex", "age"), sep = 1)
```

```
## Warning: package 'bindrcpp' was built under R version 3.4.3
```

```
## # A tibble: 76,046 x 6
##   country      year var  sex  age  value
##   * <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp   m    014     0
## 2 Afghanistan 1998 sp   m    014    30
## 3 Afghanistan 1999 sp   m    014     8
## 4 Afghanistan 2000 sp   m    014    52
## 5 Afghanistan 2001 sp   m    014   129
## 6 Afghanistan 2002 sp   m    014    90
## 7 Afghanistan 2003 sp   m    014   127
## 8 Afghanistan 2004 sp   m    014   139
## 9 Afghanistan 2005 sp   m    014   151
## 10 Afghanistan 2006 sp   m    014   193
## # ... with 76,036 more rows
```

```
who1 <- who %>%
  gather(new_sp_m014:newrel_f65, key = "key", value = "cases", na.rm = TRUE)
glimpse(who1)
```

```
## Observations: 76,046
## Variables: 6
## $ country <chr> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanis...
## $ iso2 <chr> "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", ...
## $ iso3 <chr> "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG"...
```

```
## $ year    <int> 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, ...
## $ key     <chr> "new_sp_m014", "new_sp_m014", "new_sp_m014", "new_sp_m..."
## $ cases   <int> 0, 30, 8, 52, 129, 90, 127, 139, 151, 193, 186, 187, 2...
```

```
who2 <- who1 %>%
  mutate(key = stringr::str_replace(key, "newrel", "new_rel"))
who3 <- who2 %>%
  separate(key, c("new", "type", "sexage"), sep = "_")
who3
```

```
## # A tibble: 76,046 x 8
##   country    iso2 iso3  year new  type sexage cases
## * <chr>      <chr> <chr> <int> <chr> <chr> <chr> <int>
## 1 Afghanistan AF    AFG   1997 new  sp   m014      0
## 2 Afghanistan AF    AFG   1998 new  sp   m014     30
## 3 Afghanistan AF    AFG   1999 new  sp   m014      8
## 4 Afghanistan AF    AFG   2000 new  sp   m014     52
## 5 Afghanistan AF    AFG   2001 new  sp   m014    129
## 6 Afghanistan AF    AFG   2002 new  sp   m014     90
## 7 Afghanistan AF    AFG   2003 new  sp   m014    127
## 8 Afghanistan AF    AFG   2004 new  sp   m014    139
## 9 Afghanistan AF    AFG   2005 new  sp   m014    151
## 10 Afghanistan AF    AFG   2006 new  sp   m014    193
## # ... with 76,036 more rows
```

```
who3 %>%
  count(new)
```

```
## # A tibble: 1 x 2
##   new      n
##   <chr> <int>
## 1 new   76046
```

```
who4 <- who3 %>%
  select(-new, -iso2, -iso3)
who5 <- who4 %>%
  separate(sexage, c("sex", "age"), sep = 1)
who5
```

```
## # A tibble: 76,046 x 6
##   country    year type sex  age  cases
## * <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp    m    014      0
## 2 Afghanistan 1998 sp    m    014     30
## 3 Afghanistan 1999 sp    m    014      8
## 4 Afghanistan 2000 sp    m    014     52
## 5 Afghanistan 2001 sp    m    014    129
## 6 Afghanistan 2002 sp    m    014     90
## 7 Afghanistan 2003 sp    m    014    127
## 8 Afghanistan 2004 sp    m    014    139
## 9 Afghanistan 2005 sp    m    014    151
## 10 Afghanistan 2006 sp    m    014    193
## # ... with 76,036 more rows
```

1.

In this case study I set `na.rm = TRUE` just to make it easier to check that we had the correct values. Is this reasonable? Think about how missing values are represented in this dataset. Are there implicit missing values? What's the difference between an NA and zero? More information must be known in order to know more about the data generating process. There are zeros in the data, which may indicate no cases.

```
who1 %>%  
  filter(cases == 0) %>%  
  nrow()
```

```
## [1] 11080
```

It appears that either a country has all of its values as non-missing if the World Health Organization collected data for that country or all of its values are truly non-missing. Therefore, it is okay to treat explicitly and implicitly missing values equally, and we do not lose any information by dropping the missing values.

```
gather(who, new_sp_m014:newrel_f65, key = "key", value = "cases") %>%  
  group_by(country, year) %>%  
  mutate(missing = is.na(cases)) %>%  
  select(country, year, missing) %>%  
  distinct() %>%  
  group_by(country, year) %>%  
  filter(n() > 1)
```

```
## # A tibble: 6,968 x 3  
## # Groups:   country, year [3,484]  
##   country      year missing  
##   <chr>      <int> <lgl>  
## 1 Afghanistan 1997 F  
## 2 Afghanistan 1998 F  
## 3 Afghanistan 1999 F  
## 4 Afghanistan 2000 F  
## 5 Afghanistan 2001 F  
## 6 Afghanistan 2002 F  
## 7 Afghanistan 2003 F  
## 8 Afghanistan 2004 F  
## 9 Afghanistan 2005 F  
## 10 Afghanistan 2006 F  
## # ... with 6,958 more rows
```

2.

What happens if you neglect the `mutate()` step? (`mutate(key = stringr::str_replace(key, "newrel", "new_rel"))`) If you neglect the `mutate()` step, `separate` will give the warning message “too few values”. If we check the rows for keys beginning with “newrel_”, `sexage` is missing and `type` is equal to `m014`.

```
who3nomut <- who1 %>%  
  separate(key, c("new", "type", "sexage"), sep = "_")
```

```
## Warning: Too few values at 2580 locations: 73467, 73468, 73469, 73470,  
## 73471, 73472, 73473, 73474, 73475, 73476, 73477, 73478, 73479, 73480,  
## 73481, 73482, 73483, 73484, 73485, 73486, ...
```

```
filter(who3nomut, new == "newrel") %>% head()
```

```
## # A tibble: 6 x 8
```

	country	iso2	iso3	year	new	type	sexage	cases
	<chr>	<chr>	<chr>	<int>	<chr>	<chr>	<chr>	<int>
## 1	Afghanistan	AF	AFG	2013	newrel	m014	<NA>	1705
## 2	Albania	AL	ALB	2013	newrel	m014	<NA>	14
## 3	Algeria	DZ	DZA	2013	newrel	m014	<NA>	25
## 4	Andorra	AD	AND	2013	newrel	m014	<NA>	0
## 5	Angola	AO	AGO	2013	newrel	m014	<NA>	486
## 6	Anguilla	AI	AIA	2013	newrel	m014	<NA>	0

3.

I claimed that `iso2` and `iso3` were redundant with `country`. Confirm this claim. Based on the output, `iso2` and `iso3` were redundant with `country`.

```
select(who3, country, iso2, iso3) %>%
  distinct() %>%
  group_by(country) %>%
  filter(n() > 1)
```

```
## # A tibble: 0 x 3
## # Groups:   country [0]
## # ... with 3 variables: country <chr>, iso2 <chr>, iso3 <chr>
```

4.

For each country, year, and sex, compute the total number of cases of TB. Make an informative visualisation of the data.

```
who5 %>%
  group_by(country, year, sex) %>%
  filter(year > 1995) %>%
  summarize(cases = sum(cases)) %>%
  unite(country_sex, country, sex, remove = FALSE) %>%
  ggplot(aes(x = year, y = cases, group = country_sex, colour = sex)) + geom_line()
```